

Interviewer Upskilling Simulation

A mock interview in which the AI plays the candidate and a hiring manager practises interviewing. The manager is the one being assessed — against a rubric, not against a job description. The output is an analytical estimate of how they interviewed — no pass, no fail, no gate.

CLIENT	PRODUCT	STATUS	IMMEDIATE ASK
Atibhee · Airtel	SkillBrew.AI	Draft for review	Small POC, ~2 days

01 What changed, and why

Version 2 of this document described a platform that scores *candidates* on technical depth, system design and code quality. That is the wrong product for this client. Atibhee's objective is to **upskill hiring managers**. The candidate is simulated; the manager is the learner; the report is about the manager.

Two consequences follow, and they are the substance of this revision.


CHANGE 1 – THE ASSESSED SUBJECT FLIPS

The rubric no longer measures whether the candidate is any good. It measures whether the **manager interviewed well**: five fixed criteria, owned by Atibhee's L&D function, identical across every role and every session. Comparability across 3,000 managers depends on that rubric being fixed.

CHANGE 2 – THE JOB DESCRIPTION STOPS DRIVING ANYTHING

Previously a JD was parsed into skills, and those skills generated the expectations and the persona. That coupling is removed. A role is now a **job title and a two-to-three line description**, nothing more. It gives the persona a plausible background and gives the manager something to explain — but it never generates the rubric.

What replaces the JD as the organising input is **the competency the manager needs to practise**. The trainer picks a weakness; the system picks a candidate persona that forces that weakness into the open.



The input that drives the session changes. In v2 the job description cascaded all the way to the score. In v3 the entry point is the manager competency being trained; the role family and the job card sit to one side, supplying realism only. The rubric is constant, which is what makes one manager's result comparable to another's.

02 Business context

Objective	Upskill hiring managers inside the organisation.
Method	Mock AI interview simulation. AI plays the candidate; the hiring manager interviews it.
Primary audience	Frontline managers hiring frontline workers — initially network and sales.
Manager backgrounds	Network, marketing, sales, engineering, finance. Telecom sector.
Scale	500 interviews / month at launch · 3,000+ managers in total
Commercials	₹250 per 15–20 min · ₹450 per 30 min · charged per <i>completed</i> interview
Licensing	Credit-based wallet. Creating roles and simulations is free; only completed interviews draw credit.
Integration	To sit alongside an existing online learning module.
Collaboration	Atibhee's team co-authors the simulations for each frontline profile.

COMMERCIAL CONSEQUENCE – DEFINE "COMPLETED"

Revenue is per *completed* interview, so "completed" is a billing term and needs a contractual definition before launch. A session abandoned at minute two, one ended by a network drop, and one cut short by the manager at minute twelve are three different cases. Proposal: a session is billable when it reaches the wrap-up stage or exceeds 60% of its booked duration; anything less is refunded to the wallet automatically.

03 What replaces the job description

Two objects, both deliberately thin.

3.1 Role family

A closed list. It supplies vocabulary, plausible work history, and the register a candidate of that type would speak in. It carries no assessment weight.

FAMILY	SUPPLIES	TYPICAL FRONTLINE EXAMPLE
sales	Targets, territory, channel partners, retention talk	Field sales executive, retail store promoter
marketing	Campaigns, brand, local activation, agency vocabulary	Regional marketing coordinator
technical	Field equipment, tickets, uptime, escalation vocabulary	Network field engineer, tower technician

RECOMMENDATION

Atibhee names **network and sales** as the launch frontline profiles. "Technical" as written above leans engineering; a telecom network field role is a distinct register. Suggest either renaming the third family **Network / Field Operations**, or adding it as a fourth. This is a naming decision for Atibhee's team, and it is cheap to change before the simulations are authored.

3.2 Job card

Replaces the job description entirely. One title, two to three lines, and the handful of facts a manager is expected to be able to explain — which is what Criterion 1 measures. The card exists so that "Hiring with Clarity" has something to be clear *about*.

FIELD	REQUIRED	PURPOSE
job_title	yes	Names the role. One line.
summary	yes	Two to three lines. Everything the persona knows about the job it applied for.
role_family	yes	One of the closed list above.
clarity_facts	optional	Day-in-the-life, reporting line, targets, shift, travel, location, growth path, comp band, next steps. Each fact the manager states is a tick on the Criterion 1 coverage checklist. Absent facts are simply not scored.

DESIGN NOTE

`clarity_facts` is the one place a role definition touches the score — and only as a *checklist of what was covered*, never as subject-matter difficulty. It is a coverage measure, not a knowledge test. No model reads the job card to decide what good looks like.

04 The assessment rubric

Five criteria, fixed, organisation-owned. Weights below are a starting proposal — **Atibhee's team should own them**, and the system treats them as configuration, not as code.

No criterion has a hard limit. Nothing caps, fails or overrides the overall result. Each criterion is judged on its own evidence and contributes at its weight, and the report reads as an analytical estimation of how the manager interviewed — closer to a diagnostic than to a grade.

Each criterion splits in two. **Deterministic signals** are counted by code from the transcript and session telemetry: they are exact, reproducible, and auditable. **Judge assessment** is the part that needs reading comprehension. Keeping the split explicit is what makes a score defensible to a manager who disagrees with it.

#	CRITERION	WEIGHT	DETERMINISTIC SIGNALS	JUDGE ASSESSES
1	Hiring with Clarity	20%	Coverage checklist — role scope, day-in-the-life, team & reporting, targets/KPIs, shift & travel, location, growth path, comp range, next steps. Candidate-question answer rate.	Did the explanation go beyond reading the job card? Was it tailored to <i>this</i> candidate?
2	Structured Interviewing	25%	Open vs closed question ratio · behavioural / STAR-eliciting question count · probing follow-up count · probe handling on vague answers and inflated claims · same-competency question consistency.	Quality of STAR completion — Situation, Task, Action, <i>Result</i> ; most managers never reach Result. Whether probes were purposeful or interrogative.
3	Unconscious Bias	20%	Protected-topic detector on manager turns: marital status, children & family planning, pregnancy, age, religion, caste or community, native place, gender-role assumptions, disability, "will your family allow", "adjust". Handling of bias-bait probes.	Subtler bias — differential warmth, assumption-loaded framing, stereotype inference from name or region.
4	Candidate Experience	20%	Greeting · self-introduction · agenda and duration stated · time to first invitation for candidate questions · abrupt-end detector (no close, next steps or thanks) · manager interruptions · silence handling.	Warmth. Whether the candidate could plausibly leave with a good impression of the employer.
5	Communication & Tone	15%	Talk-to-listen ratio · longest monologue · words per minute · filler density · interruptions · aggression / curtness classifier per manager turn · composure during the provocation probe.	Tone trajectory across the call, clarity for a frontline (non-corporate) audience, jargon load.

NO HARD LIMITS – THIS IS AN ESTIMATE, NOT A GATE

No criterion caps, fails or overrides the result. Every criterion is scored on its own evidence and contributes at its weight; the output is an **analytical estimation of how this manager interviewed**, not a pass/fail certification.

Unconscious Bias is therefore reported like the others — a score, the exact turns that produced it, and what to do differently. It is the criterion most likely to matter to the organisation, so the report surfaces it prominently and quotes it precisely, but prominence is a presentation decision, not a scoring rule.

05 Persona configuration by traits

A persona is a bundle of trait values plus a set of behaviours. It is **independent of the role family** — the same "Vague Responder" works as a sales candidate or a network technician, because vagueness is a behaviour, not a subject.

5.1 Trait axes

Every persona is a point in this space. All axes are integers 0–10.

AXIS	WHAT IT CHANGES IN THE ROOM
smartness / dumbness	Ceiling of insight. The ratio between them is what the manager is trying to read.
seriousness	How much the candidate treats this as mattering.
effort	Whether they work to answer, or wait to be dragged.
interest	Engagement with this role specifically. Low interest tests whether the manager can tell disinterest from inability.
honesty	Gap between what is claimed and what is true. Drives the inflated-claim probe.
preparedness	Whether they researched the company and the role.
nervousness	Drives hesitation, self-correction and pauses. High values test warmth.

DELIVERY LAYER, CONFIGURED ALONGSIDE THE TRAITS

filler_frequency and hesitation_frequency (0–10), pace, verbosity, formality, verbal tics, whether the candidate may interrupt, and a per-skill knowledge_ceiling (0–10) that the persona can never exceed however hard the manager pushes.

5.2 Persona library — each one exists to stress a criterion

This is the heart of the change. A trainer does not pick a persona because it is realistic; they pick it because it **forces a specific manager competency into the open**. Every persona in the library must earn its place by stressing something the rubric measures.

PERSONA	TRAIT SIGNATURE	STRESSES	WHAT THE MANAGER HAS TO DO
Vague Responder	effort 3 · smartness 4 · honesty 6	C2	Answers in generalities. Only concrete detail reaches a STAR Result — the manager must probe three times or leave with nothing.
Inflated CV	honesty 3 · preparedness 7 · smartness 5	C2	Claims ownership of work they supported. The manager must convert "we" to "I" without accusing.
Curt / Unprofessional	seriousness 3 · interest 4 · may_interrupt	C5 C4	Short, faintly dismissive, interrupts. Tests composure — does the manager mirror the tone or hold it?
Off-Topic Asker	preparedness 3 · interest 7	C1 C4	Asks irrelevant or premature questions (salary in minute two). Tests redirection without dismissal, and whether role facts get covered anyway.
Bias Bait	honesty 8 · nervousness 6	C3	Volunteers protected-category information unprompted — marriage plans, family, native place. Tests whether the manager follows the thread or steers away. The single most important persona in the library.
Nervous but Capable	nervousness 8 · smartness 7 · effort 7	C4	Real ability behind visible anxiety. Rewards warmth; punishes managers who read nerves as incompetence.
Rambler	verbosity high · seriousness 6	C5 C2	Long, drifting answers. Tests time control and courteous interruption inside a fixed slot.
Monosyllabic	effort 2 · nervousness 5 · interest 4	C4 C2	Yes / no answers and silences. Tests whether the manager can open a closed room, and how they handle dead air.

WHY THIS IS THE RIGHT SHAPE

Because the rubric is fixed and the personas are role-agnostic, coverage becomes a grid: **8 personas × 3 role families = 24 simulations** from one authoring effort, with no job descriptions written at all. Atibhee's team tunes trait values and adds personas; neither requires engineering work, which is also what makes "no charge for creating roles and simulations" economically safe.

OPEN QUESTION FOR ATIBHEE

Does the trainer choose the persona, or does the system choose it from the manager's weakest prior criterion? Manager-chosen is simpler and ships in the POC. System-chosen is the stronger learning product — it closes the loop by sending each manager the candidate they handle worst — but it needs history, so it belongs after launch.

06 Session flow

Mapped directly to the flow Atibhee described.

#	STAGE	OBSERVED FOR	NOTES
1	Manager greets and introduces themselves	C4	Greeting, self-introduction, agenda and duration are each a deterministic tick.
2	AI candidate introduces itself	—	Spoken from a fixed opening line so every session starts in character.
3	Manager explains the role	C1	Coverage checklist runs here and stays open for the rest of the call.
4	Behavioural questions and probing	C2	The main body. Persona behaviour determines how hard probing has to work.
5	Scripted disruption	C5 C3	The persona introduces its designed difficulty — curtness, an off-topic question, or bias bait — at a controlled point rather than at random, so sessions stay comparable.
6	Candidate questions	C1 C4	Answer rate is counted; quality is judged.
7	Close and next steps	C4	Abrupt-end detector fires if the call ends with no close, next steps or thanks.

DESIGN NOTE – DISRUPTIONS ARE PLACED, NOT RANDOM

Atibhee asked for scenarios like unprofessional tone and irrelevant questions. If those fire randomly, two managers get different difficulty and their scores stop being comparable. Each persona therefore carries a *scripted* disruption at a defined point in the session. Same persona, same challenge, every manager.

07 The feedback report

Atibhee asked for three things: what the manager did perfectly, where they need to focus, and areas for development. The report is structured exactly that way, and **every statement in it is anchored to a quoted turn** with a timestamp. Unevidenced feedback is not actionable and will not survive contact with 3,000 managers.

SECTION	CONTENTS
Outcome	Weighted estimate across the five criteria, with a confidence note where the session was too short or too quiet to judge a criterion fairly. No pass/fail verdict.
Criterion breakdown	Per criterion: score, the deterministic counts behind it, and the judge's reasoning. The counts are shown because they are checkable.
Done well	Two to four specific moments, quoted.
Focus areas	The lowest-scoring criteria, each with the moment it went wrong and what to have said instead.
Development plan	Which persona to practise against next, and why.
Transcript	Full, timestamped, with the manager's and the candidate's turns separated. Doubles as the audit record.

AGGREGATE REPORTING – WORTH NAMING NOW

With 3,000 managers, the organisational view is likely to be the thing Atibhee's leadership actually buys: bias-flag rate by business unit, STAR completion trend, average talk-to-listen ratio. It is out of scope for the POC, but the report schema should be designed so it can be aggregated later rather than retrofitted.

08 First runnable milestone

Before any client-facing scope is fixed, the product has to be real enough to sit down with: **pick a persona, hold an actual interview with it, and read the report afterwards**. Billing, accounts, LMS integration and aggregate reporting are all downstream of that, and none of them is worth scoping until the conversation itself is convincing.

NEEDED TO CONVERSE AND ASSESS	DELIBERATELY DEFERRED
• The full persona library, any one selectable at session start • One job card per role family — title and three lines • Live two-way session of realistic length • Turn-by-turn transcript of both sides • Deterministic signals counted from the transcript • Judge pass across all five criteria • Analytical report with quoted evidence	• Credit wallet and billing • LMS integration • Manager accounts, history, dashboards • Aggregate or leadership reporting • System-chosen persona selection • Client-facing POC packaging

WHY THIS ORDERING

A persona can read convincingly on paper and fall apart in conversation, and no amount of specification settles which. Interviewing every persona in the library is the only way to find out whether each one actually stresses the criterion it was designed to stress — and that judgement has to come before scope is promised to anyone.

09 Decisions needed from Atibhee

REF	QUESTION	WHY IT MATTERS
D-1	Confirm or replace the five weights.	Stated as a starting proposal. The rubric is Atibhee's instrument, and weights encode what the organisation values.
D-2	What score bands should the estimate map to, and what are they called?	Decided: no criterion gates the result. Remaining question is presentation — bands need names Atibhee's managers will accept as developmental rather than punitive.
D-3	Definition of a "completed" interview.	It is the billing unit. Needs to be contractual, not implied.
D-4	Third role family: Technical, or Network / Field Operations?	Launch profiles are network and sales. Cheap now, expensive after simulations are authored.
D-5	Who picks the persona — trainer or system?	Manager-chosen ships in the POC; system-chosen is the stronger learning loop.
D-6	Language and code-mixing.	Frontline interviews in Indian telecom are frequently conducted in Hindi or mixed Hindi-English. This materially affects the voice model, the transcript and every deterministic signal. Not yet scoped.
D-7	Recording consent and retention.	Sessions record a real manager's voice. Needs a policy before the first live user.
D-8	Does a manager's report go to their reporting line?	Determines whether this is a safe practice space or an evaluation. It changes manager behaviour inside the session, and therefore what the scores mean.

D-6 AND D-8 DESERVE EARLY ATTENTION

Language is the larger engineering risk in this document: a Hindi or code-mixed session breaks talk-time analysis, filler detection, the protected-topic detector and STAR extraction all at once. If frontline interviews are not conducted in English, that must surface now rather than after the POC.

D-8 is the larger product risk. A manager who believes their score reaches their boss will interview defensively, and the tool measures performance rather than habit.